

어텐션 매커니즘 for 스터디

Tae Hwan Jung

1. Seq2Seq

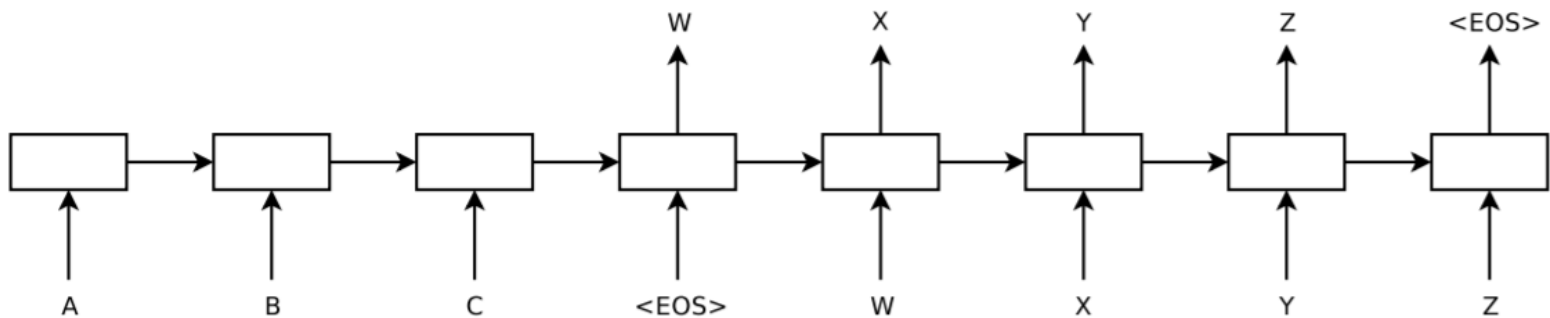


Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

1. Seq2Seq

Encoder

Decoder




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

1. Seq2Seq

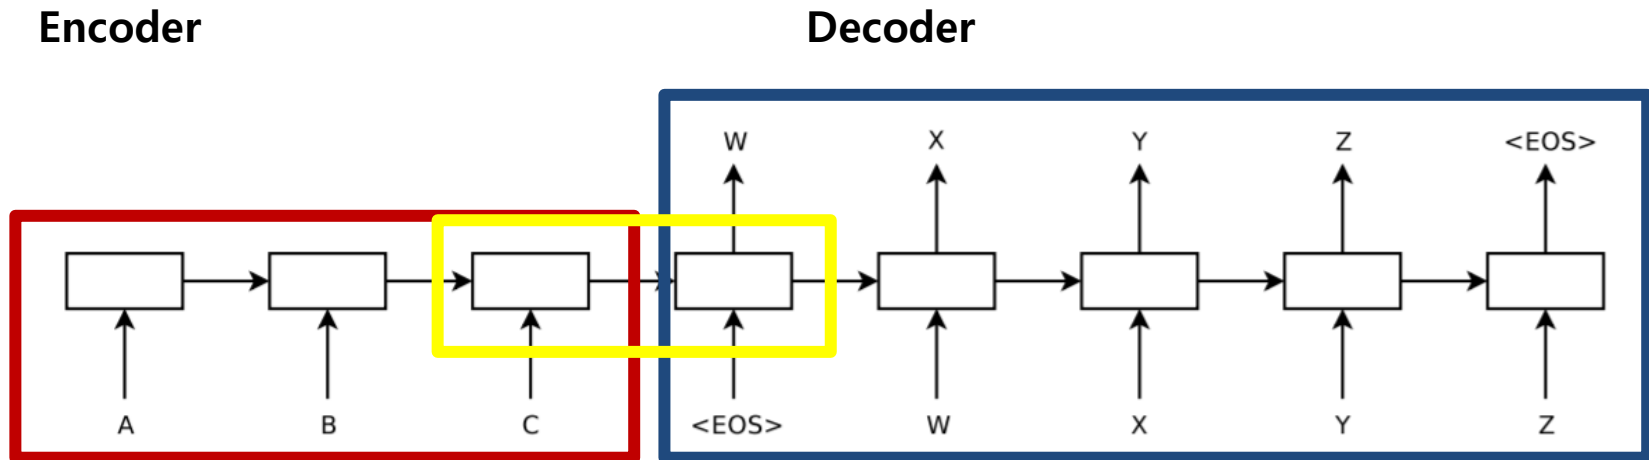


Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

Encoder의 최종 hidden State가 Decoder의 initialize Hidden State로 들어간다.

- <https://arxiv.org/pdf/1409.3215.pdf> 출처

1. Seq2Seq

Inputs: input, h_0

- **input** of shape $(seq_len, batch, input_size)$: tensor containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` or `torch.nn.utils.rnn.pack_sequence()` for details.
- **h_0** of shape $(num_layers * num_directions, batch, hidden_size)$: tensor containing the initial hidden state for each element in the batch. Defaults to zero if not provided. If the RNN is bidirectional, `num_directions` should be 2, else it should be 1.

1. Seq2Seq

Outputs: output, h_n

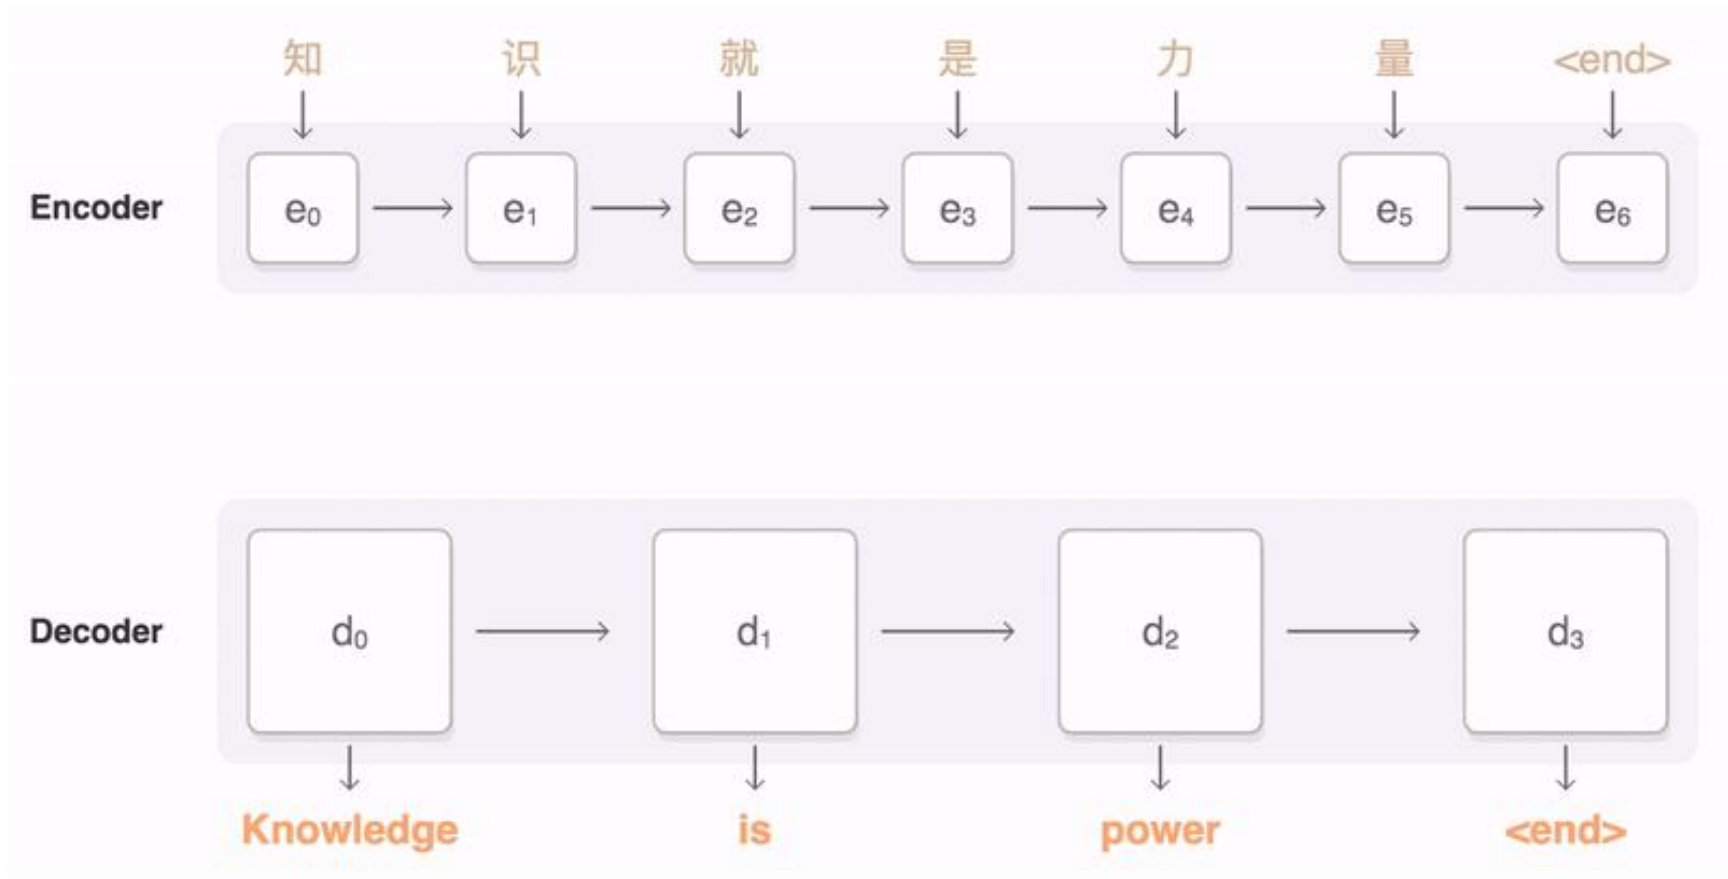
- **output** of shape $(seq_len, batch, num_directions * hidden_size)$: tensor containing the output features (h_k) from the last layer of the RNN, for each k . If a `torch.nn.utils.rnn.PackedSequence` has been given as the input, the output will also be a packed sequence.

For the unpacked case, the directions can be separated using `output.view(seq_len, batch, num_directions, hidden_size)`, with forward and backward being direction 0 and 1 respectively. Similarly, the directions can be separated in the packed case.

- **h_n** $(num_layers * num_directions, batch, hidden_size)$: tensor containing the hidden state for $k = seq_len$.

Like *output*, the layers can be separated using `h_n.view(num_layers, num_directions, batch, hidden_size)`.

2. Seq2Seq-Attention



- <https://google.github.io/seq2seq/> 출처

2. Seq2Seq-Attention




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

$$e_{ij} = a(s_{i-1}, h_j)$$

E_ij = alignment model(디코더의 이전 hidden state, 인코더의 hidden state)

2. Seq2Seq-Attention

$$\vec{\alpha}_i = [\alpha_{i1}, \alpha_{i2}, \dots, \alpha_{iT_x}]$$




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})}$$

2. Seq2Seq-Attention

$$\vec{c}_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j = \boxed{F} \alpha_i$$




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

2. Seq2Seq-Attention

```
model[i] = self.out(torch.cat((dec_output, context), 1))
```




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

2. Seq2Seq-Attention

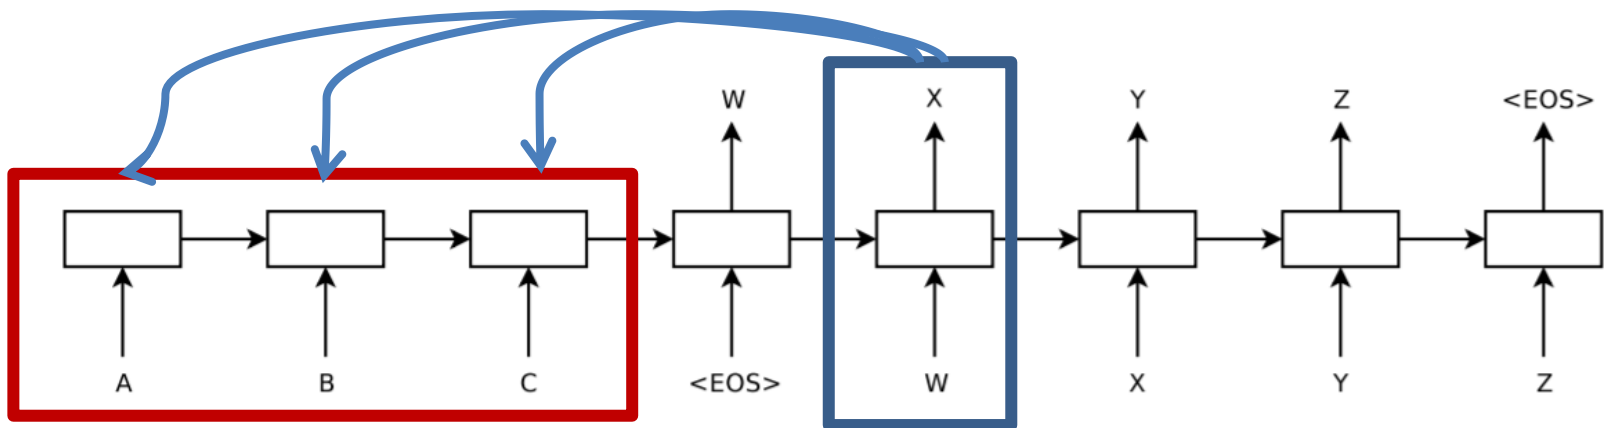


Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

2. Seq2Seq-Attention




Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

2. Seq2Seq-Attention

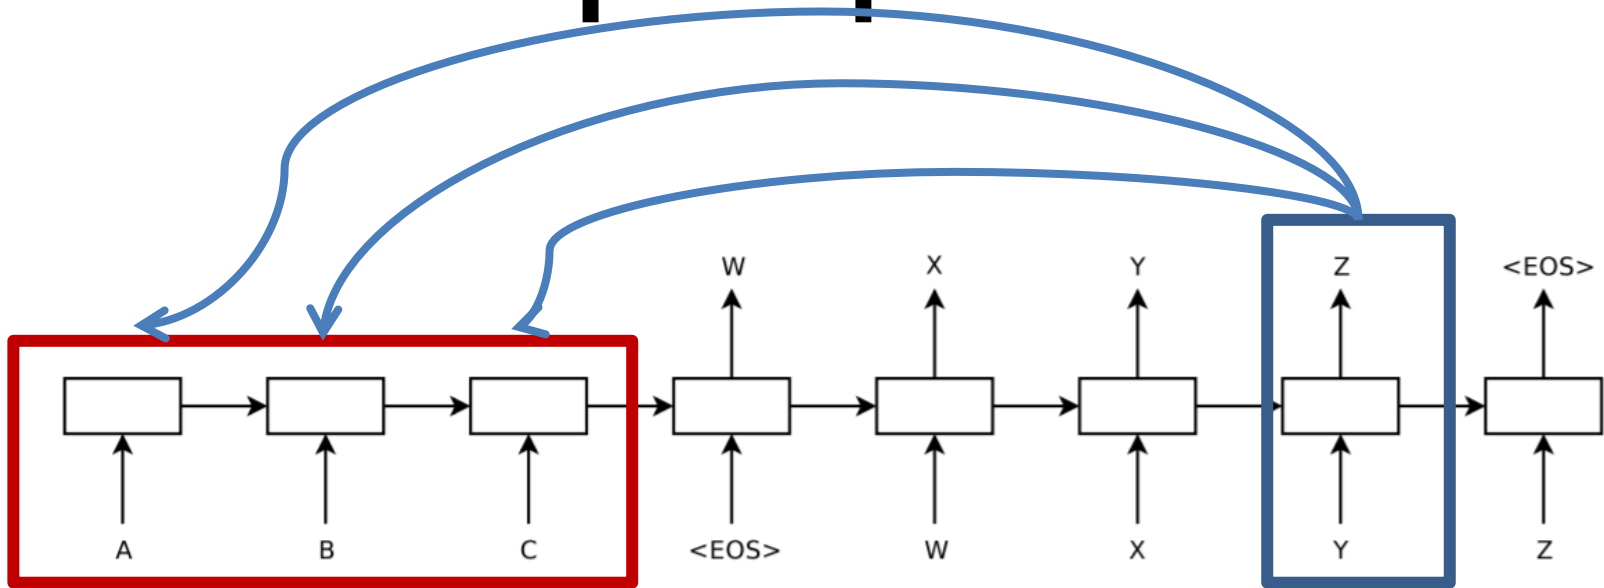


Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.

2. Seq2Seq-Attention

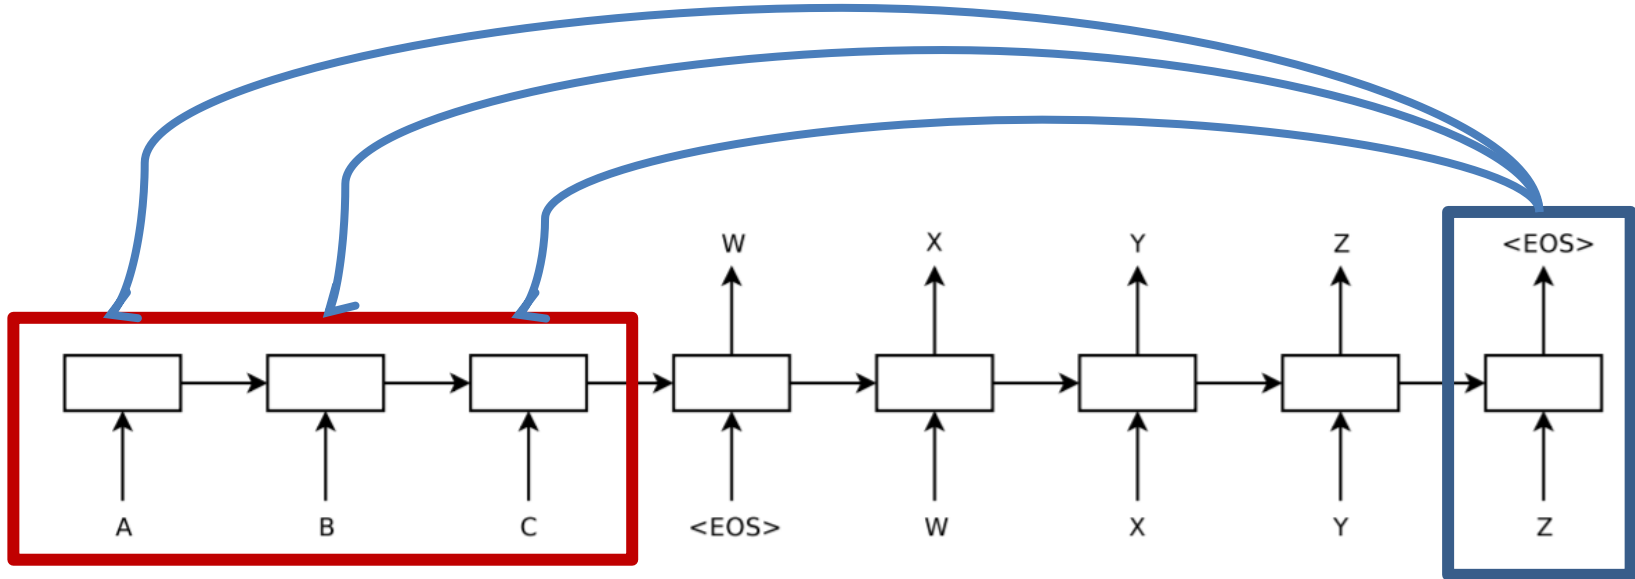


Figure 1: Our model reads an input sentence “ABC” and produces “WXYZ” as the output sentence. The model stops making predictions after outputting the end-of-sentence token. Note that the LSTM reads the input sentence in reverse, because doing so introduces many short term dependencies in the data that make the optimization problem much easier.